

# Draft

Diana Lin & Nima Jamshidi

14/03/2020

## Contents

0.1	Introduction . . . . .	1
0.2	Data Description . . . . .	1
0.3	Exploring the Dataset . . . . .	2
0.4	Methods . . . . .	6
0.5	Results . . . . .	7
0.6	Discussion . . . . .	7
0.7	Conclusion . . . . .	7

## 0.1 Introduction

The dataset we have chosen to work with is the “Medical Expenses” dataset used in the book Machine Learning with R, by Brett Lantz. This dataset was extracted from Kaggle by Github user @meperezcuello. The information about this dataset has been extracted from their GitHub Gist.

This dataset is very interesting as the USA does not have universal healthcare, and is known for bankrupting its citizens with hospital visits despite having insurance. It will be interesting to see the relationship between characteristics of a beneficiary, such as BMI and Smoking status, and the **charges** incurred.

Originally, this dataset was used to train a machine learning algorithm to accurately predict insurance costs using linear regression.

## 0.2 Data Description

This dataset explains the medical insurance costs of a small sample of the USA population. Each row corresponds to a beneficiary. Various metadata was recorded as well.

```
# import the data
costs <- read_csv(
  here("data", "raw", "data.csv"),
  col_types = cols(
    age = col_integer(),
    sex = readr::col_factor(),
    bmi = col_double(),
    children = col_integer(),
    smoker = readr::col_factor(),
```

```

    region = readr::col_factor(),
    charges = col_double()
  )
)

```

The columns (except the last one) in this dataset correspond to metadata, where the last column is the monetary charges of medical insurance:

```
colnames(costs)
```

```
## [1] "age"      "sex"      "bmi"      "children" "smoker"   "region"   "charges"
```

Here are the possible values for each of the above column names:

Variable	Type	Description
Age	integer	the primary beneficiary's age in years
Sex	factor	the beneficiary's sex: <b>female</b> or <b>male</b>
BMI	double	the beneficiary's Body Mass Index, a measure of their body fat based on height and weight (measured in kg/m2), an ideal range of 18.5 to 24.9
Children	integer	the number of dependents on the primary beneficiary's insurance policy
Smoker	factor	whether or not the beneficiary is a smoker: <b>yes</b> or <b>no</b>
Region	factor	the beneficiary's residential area in the USA: <b>southwest</b> , <b>southeast</b> , <b>northwest</b> , or <b>northeast</b>
Charges	double	the monetary charges the beneficiary was billed by health insurance

### 0.3 Exploring the Dataset

Here is a summary of the dataset, and the values of each variable:


```
summary(costs)
```

```
##      age      sex      bmi      children      smoker
##  Min.   :18.00  female:662  Min.    :15.96  Min.     :0.000  yes: 274
##  1st Qu.:27.00  male  :676  1st Qu.:26.30  1st Qu.:0.000  no :1064
##  Median :39.00
##  Mean   :39.21
##  3rd Qu.:51.00
##  Max.   :64.00
##      region      charges
## southwest:325  Min.    : 1122
## southeast:364  1st Qu.: 4740
## northwest:325  Median   : 9382
## northeast:324  Mean     :13270
##               3rd Qu.:16640
```

## Max. :63770

### 0.3.1 Correlogram


In this section we are inspecting the data set to see if there is any correlation between the variables. From now on we want to consider charges as our dependent variable. In order to analyze correlation between variables, the ones that are categorical with two categories, are translated into binary vectors. The only categorical variable with more than two categories, is region. We split this variable into four different binary vectors, each indicating if the sample data has category (1) or not (0). After using dummy variables for sex, smoker, and region, according to the correlogram below, smoker and charges has the strongest correlation of 0.79. No high collinearity between independent variables is observed.



### 0.3.2 Faceted Plot


Here we want to explore the data to see if there is any cluster of data points. While the data between regions and sex does not appear to vary much, the smokers vs nonsmokers of each facet appear to cluster together, with the non-smokers having an overall lower medical cost.

### Exploring the Medical Costs Dataset



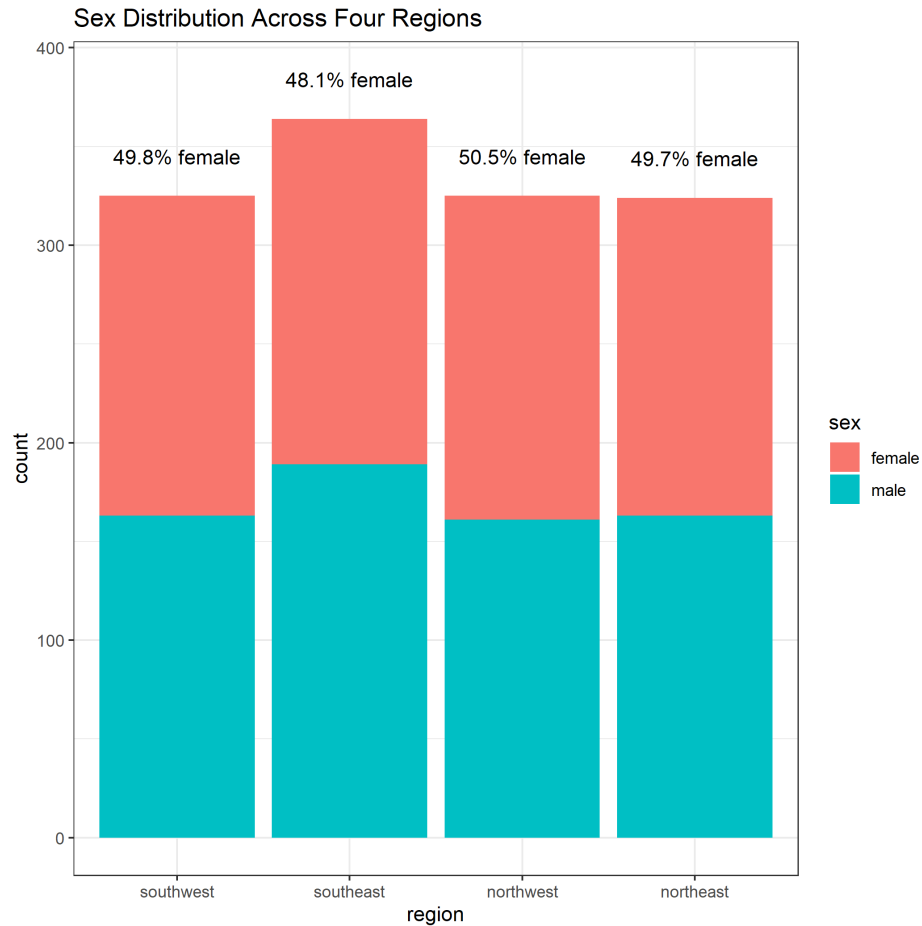
### 0.3.3 Histogram

How is the distribution of sex among different age groups? Looking at the dataset, there appears to be more beneficiaries in the 20-60 age range. The biggest difference in the number of beneficiaries from different sex is seen in the 20-30 bracket.



#### 0.3.4 Stacked Bar Chart

How about the distribution of sex among the regions? This plot shows the distribution of sex in each of the four regions. At a glance, the dataset looks very even when it comes to sex, but there are slightly more beneficiaries in the southeast.



## 0.4 Methods

Here we use multiple linear regression to study the relations between the independent variables and the dependent one, charges. Below you can find the results of the regression in Table 2. 'lm' function in R transforms a categorical variable with n levels into n-1 variables each with two levels to insure the variables are independent. Here we can see that variables age, bmi, children, and smoker are significantly important in the regression. Sex is an insignificant factor in the model.

```
tidied <- readRDS("../data/linear_model/tidied.rds")
kable(tidied, caption = "Summary of the model's variables and their respective coefficients")
```

In Table ?? we can see that the r-squared value is "r glanced\$r.squared".

```
glanced <- readRDS("../data/linear_model/glanced.rds")
kable(glanced)
```

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residu
0.750913	0.7494136	6062.102	500.8107	0	9	-13547.75	27115.51	27167.5	48839532844	13

Table 2: Summary of the model's variables and their respective coefficients

term	estimate	std.error	statistic	p.value
(Intercept)	-11938.5386	987.81918	-12.0857530	0.0000000
age	256.8564	11.89885	21.5866552	0.0000000
sexmale	-131.3144	332.94544	-0.3944020	0.6933475
bmi	339.1935	28.59947	11.8601306	0.0000000
children	475.5005	137.80409	3.4505546	0.0005770
smokeryes	23848.5345	413.15335	57.7232020	0.0000000
regionnorthwest	-352.9639	476.27579	-0.7410914	0.4587689
regionsoutheast	-1035.0220	478.69221	-2.1621870	0.0307817
regionsouthwest	-960.0510	477.93302	-2.0087563	0.0447649

## 0.5 Results

*# PLACEHOLDER FOR LINEAR REGRESSION*

## 0.6 Discussion

*# PLACEHOLDER FOR LINEAR REGRESSION*

## 0.7 Conclusion

*# PLACEHOLDER FOR LINEAR REGRESSION*